Dongle: A lightweight interface between EI, EA and AS processes and the Presentation Process

Russell G Almond

`ralmond@fsu.edu`

1 Introduction

In the for process architecture (Almond, Steinberg, & Mislevy, 2002), the presentation process often has real-time constraints. In an adaptive system, it is unacceptable for the subject to wait for more than a few moments for the presentation process to load the next task. Similarly, information about subject performance should be returned quickly to the presentation process to avoid annoying delays. However, there is considerable uncertainty about how long the evidence identification (EI), evidence accumulation (EA) and activity selection (AS) will take to run.

The "Dongle" is a lightweight adapter between the presentation process and the other three processes consisting of a database, a web server and a number of PHP scripts. When the presentation process needs information from one of the other processes, it sends an HTTP POST message to the appropriate script. The script fetches the appropriate record for the given user and application from the database and returns it as a JSON document. The other processes, when they complete their work, update the records in the database. If the EI, EA or AS processing is not completed when the presentation process requests the information, the most recent available information from the database is provided. This might be slightly out of date, but this is often better than waiting. If no record is a available for a given user, a default record for that application is returned. In this way the dongle process is as fast as the load on the web server and network latency allows.

Figure 1 shows the basic architecture. At its heart, it is a modified LAMP (Linux, Apache, Mongo®,[1] PHP stack. The game engine communicates to the other processes through HTTP requests to the web server, and the other three processes (EI—evidence identification, EA—evidence accumulation, and AS—activity selection) communicate through the database.

2 Proc4 Message Format

The Proc4 message is an object consisting of two parts: a header and a body. The header is a series of mandatory and optional fields which are used to route and prioritize the message. The data portion is a container object which can hold anything. In general, its value will be determined by the message. Listing 1 shows the data structure in JSON (java script object notation; Bassett, 2015) format.

The fields have the following definition

---

[1]The modification is that the Mongo database is substituted for the MySQL database used in many web
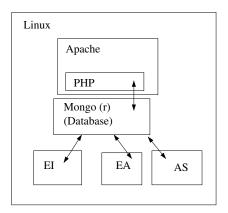
Figure 1: Proc4 Dongle Architecture.

Listing 1: A typical Proc4 message in JSON format.

```
1  {
2    app: "ecd://epls.coe.fsu.edu/PP",
3    uid: "Student 1",
4    context: "SpiderWeb",
5    sender: "Evidence Identification",
6    mess: "Task Observables",
7    timestamp: "2018-10-22 18:30:43 EDT",
8    processed: false,
9    data:{
10     trophy: "gold",
11     solved: true,
12     objects: 10,
13     agents: ["ramp","ramp","springboard"],
14     solutionTime: {time:62.25, units:"secs"}
15   }
16 }
```

app (String, Required). A globally unique identifier (guid) for the assessment application. This should have a URL-like syntax with the first header corresponding to the domain of the organization issuing the ID.

uid (String, Required). A unique identifier for the student or player. Note that administrative messages which correspond to all players could have an empty string as a value.

context (String, Required) A unique identifier for the context in which the message was generated. In *Physics Playground* this corresponds to game levels, but it might have other meanings in other applications.

sender (String, Optional). An identifier for the process which generated the message.

mess (String, Required). An identifier for the contents of the message.

timestamp (String in POSIX format, Required). A timestamp for the message. Generally, messages for the same `uid` need to be processed in chronological order.

processed (Boolean, Optional). A flag that can be set after the message has been processed.

error (Any, Optional). If an error occurs during the processing of the message, details can be placed here.

data (Any, Required). The contents of the message. The expected content and format is controlled by the application and vocabulary.

Note that this P4 Message class can be extended with additional header fields. For example, Evidence Identification messages add `verb` and `object` fields.

The `app` field plays a special role in defining the *vocabulary* for the assessment. In particular, the application defines the legal values for the `context`, `mess` and `sender` can take, as well as the expected structure of the `data` for various values of `mess`

The dongle process uses a document-orient database (Mongo®; "The MongoDB 4.0 Manual", 2018). This allows indexes to be built for the header fields, while allowing the data fields to be unrestricted. The database serves as a buffer between the game process, communicating through the web server, and EI, EA and AS processes communicating directly with the database.

---

applications.

3   Mongo Database Schemas

There are four database collections in the Proc4 database: `AuthorizedApps`, `Players`, `Statistics`, and `Activities`. The latter three provide mechanisms for the EI, EA and AS processes respectively to communicate with the game engines. When queried, all three will return the record (or the latest record) matching the provided `app` and `uid`. If no record is available with those ids, then a special default record for that application, with `uid="*DEFAULT*"`, will be returned.

The configuration resides in two places. The file `/usr/local/share/Proc4/Proc4.js` contains initialization codes. In particular, it contains both the definition of the authorized applications, and the credentials for the various databases (if secure login is enabled). The script file `setupDatabases.js` sets up the security for the four processes (in the `Proc4`, `EIRecords`, `EARecords` and `ASRecords` databases), and `setupProc4.js` sets up the collections and indexes in the `Proc4` database. Most of the code in `setupProc4.js` is described below. These files are located in the `inst/config` directory (or just `config` in the installed package) of the R `Proc4` package. They can be run with the shell command `mongo` *script*`.js`.

3.1   AuthorizedApps Collection

The `AuthorizedApps` collection has two purposes: to provide as a first line security and to provide for a graceful shutdown of the EI, EA and AS processes. The security is somewhat minimal, it is just that a message with an `app` field that does not match one of the `AuthorizedApps` records will return an error. While this will not stop a determined attacker, it will make it harder for a hacker randomly looking for open ports to get data from the database. The second is related to the `active` field of the record. The EI, EA, and AS processes periodically check this, and if it is false, then they shut down gracefully.

This collection was never give an formal schema. Listing 2 shows the code that creates it. The `app` and `active` fields should be counted as required and `doc` as recommended.

This fields is mostly manipulated with raw database commands setting the `active` field. The launch scripts for the EI and EA processes both set the `active` field to true. (They will also, if necessary, add a record for the application.) To signal a graceful shutdown, simply set the `active` field to false. Listing 3 gives some example syntaxes (to run from the mongo command shell).

The graceful shutdown causes the process to process all remaining events in the queue and then stop. It became obvious that there was a need of a rapid shutdown as well. In particular, there were cases where the processes should be shut down after finishing processing the current event so that repairs could take place before continuing. A future version of the code should have this.

Listing 2: AuthorizedApps Collection

```
1   // In Proc4.js configuration file.
2   var apps = [
3     {"app":"ecd://epls.coe.fsu.edu/P4test","doc":"Testing",
4       "active":true},
5     {"app":"ecd://epls.coe.fsu.edu/PhysicsPlayground/userControl",
6       "doc":"User␣Controlled␣(Spring␣2019)","active":true},
7     {"app":"ecd://epls.coe.fsu.edu/PhysicsPlayground/linear",
8       "doc":"linear␣(Spring␣2019)","active":true},
9     {"app":"ecd://epls.coe.fsu.edu/PhysicsPlayground/adaptive",
10      "doc":"adaptive␣(Spring␣2019)","active":true}
11  ];
12  // In setupProc4.js setup file.
13  apps.forEach(function (row) {
14    db.AuthorizedApps.replaceOne({"app":row.app},row,
15      {"upsert":true});
16  });
```

Listing 3: AuthorizedApps Collection

```
1   // Shutdown just the P4test engine
2   db.AuthorizedApps.update({app:{"$regex":"P4test"}},
3     {"$set":{active:false}});
4   // Shutdown all active apps
5   db.AuthorizedApps.update({active:true},
6     {"$set":{active:false}});
```

3.2 Players Collection

The original purpose of the `Players` collection was to notify the other processes about which players were currently active in the game, and which were not. As the game engine would send a message to the scoring server when the players logged in, the `data` field of the `Players` collection was used to store information about the player which needed to persist between game sessions; in particular, the player's bank balance and the `trophyHall`, a list of completed levels with the coin awarded for each. Listing 4 gives the schema for this collection.

Note that this is a subset of the basic `Proc4` message format. The `active` field should be set to true when the player starts play, and to false when the stop. The PHP scripts `PlayerStart.php` and `PlayerStop.php` start and stop the player. Note that current the game may or may not post a message to the `PlayerStop.php` script. Currently, the other processes are not relying on this, but this needs to be revisited later if it becomes important.

In addition to setting the active field, the `PlayerStart.php` script returns the current value of the `data` field as part of a message. Currently, `data` is a named list (or dictionary) with two elements: `bankBalance` and `trophyHall`. The first in an integer value giving the players' bank balance (as of the last event processed), and `trophyHall` is a named list with the names corresponding to levels and the values corresponding to coins ("gold" or "silver"). Levels for which a coin was not awarded do not appear in the list.

The EI process is responsible for keeping the `Players` collection up to date. There is a special collection of rules called `TrophyHallRules.json` which contain the logic for doing the update. These include special trigger rules which run to update the `Players` collection; and a special listener which listens for those messages.

3.3 Statistics Collection

The EA process, in response to each release of observables from the EI process, performs the following actions: (1) it fetches the student model for the reference `uid`, (2) it updates the student model using the new evidence, (3) it updates a list of `statistics` for the updated student model, (4) it posts the updated statistics so they can be viewed both other processes. One of the places it posts the updated statistics is in the `Statistics` collection in the `Proc4` database. Listing 5 provides the schema for that collection.

The EA process updates the `data` field of the record for the just processed student. The value of the `data` field is a named list of statistics. The value of the statistic depends on the specifications given to the EA process. There are three common kinds of statistics that are used: real-valued statistics (for example the expected a posteriori or EAP statistic), string valued statistics (for example, the mode or median of a Bayes net node), and vector valued statistics (for example, the probability distribution for a node).

The script `PlayerStats.php` returns the current record in the `Statistics` collection for the player.

Listing 4: Players Collection

```
1   db.createCollection("Players", {
2   validator: {
3       \$jsonSchema: {
4           bsonType: "object",
5           required: ["app","uid","active","timestamp"],
6           properties: {
7               app: {
8                   bsonType: "string",
9                   description: "Application␣ID␣(string)"
10              },
11              uid: {
12                  bsonType: "string",
13                  description: "User␣(student)␣ID␣(string)"
14              },
15              active: {
16                  bsonType: "bool",
17                  description: "Is␣the␣player␣currently␣active?"
18              },
19              context: {
20                  bsonType: "string",
21                  description: "Context␣(task)␣ID␣(string)"
22              },
23              timestamp: {
24                  bsonType: "date",
25                  description: "Timestamp"
26              },
27              data: {
28                  bsonType: "object",
29                  description: "Player␣State␣information␣passed␣to␣game
30              }
31          }
32      }
33  },
34  validationAction: "warn"
35 });
36 db.Players.createIndex( { app:1, uid: 1});
```

Listing 5: Statistics Collection

```
1   db.createCollection("Statistics", {
2   validator: {
3       \$jsonSchema: {
4           bsonType: "object",
5           required: ["app","uid","timestamp"],
6           properties: {
7               app: {
8                   bsonType: "string",
9                   description: "Application␣ID␣(string)"
10              },
11              uid: {
12                  bsonType: "string",
13                  description: "User␣(student)␣ID␣(string)"
14              },
15              context: {
16                  bsonType: "string",
17                  description: "Context␣(task)␣ID␣(string)"
18              },
19              sender: {
20                  bsonType: "string",
21                  description: "Who␣posted␣this␣message."
22              },
23              mess: {
24                  bsonType: "string",
25                  description: "Topic␣of␣message"
26              },
27              timestamp: {
28                  bsonType: "date",
29                  description: "Timestamp"
30              },
31              data: {
32                  bsonType: "object",
33                  description: "Named␣list␣of␣statistics."
34              }
35          }
36      }
37  },
38  validationAction: "warn"
39 });
40 db.Statistics.createIndex( { app:1, uid: 1, timestamp: -1});
```

3.4 Activities Collection

The last collection was never actually implemented, but it was created for future expansion. Listing 6 provides its draft schema. Its purpose was to provide a place the AS could post messages about which level to provide next. The script `PlayerLevel.php` would return its current value.

Although this was not implemented, the following notes describe the planned design. The data field would contain four components: `topic`, `completedLevels`, `availableLevels`, and `supportMode`. Topic is intended as a string valued field. The game levels would be divided into a number of *topics*. When the internal criteria in the AS algorithm were met, the player would "graduate" from the topic, and the value of the topic field would change. The components `completedLevels` and `availableLevels` would list all of the levels in the topic. The `availableLevels` would be sorted into the desired order. As levels were completed, they would be moved to the `completedLevels` field. The `supportMode` component is a logical variable that would be set if the player should be placed into a learning support rather than the game at the start of a new level.

The design of the `Activities` collection is designed to be robust against latency problems with the EI, EA and AS processes. In order to for the algorithm to be completely adaptive, then all three processes must complete between the time the player finishes the game level and the system requests a new game level. If this condition does not hold, the database can return the sorted list of levels, taking the just completed level from the list of available levels and putting it on the completed list. So although the next level played may not be completely optimal, as long as the EI, EA and AS processes don't fall too far behind it will be at least close to optimal.

4 PHP communication Layer

A number of PHP scripts are provided to allow the game engine (or other process) to access the information in the database. The PHP scripts always return the most recent information available for the player, or if no information is available for the player, a default record is returned. In particular, this means that the processes should never block, but they might not return the most recent information if there are still unprocessed events working their way through the EI, EA and AS processes.

All of the PHP scripts expect the headers in the basic P4 Message format using fields of an HTTP POST request. In particular, it is looking for fields of with the names "app", "uid", "context", "sender", "mess", "timestamp", and "data". If accessed using the post method, all of the pages should return a file of type 'application/json' in utf-8 encoding. If the php scripts are accessed using a GET rather than a POST request, then a HTML form with these fields is returned (to be used for testing).

There are four primary files which are used for the communication:

Listing 6: Activities Collection

```
1    db.createCollection("Activities", {
2    validator: {
3        \$jsonSchema: {
4            bsonType: "object",
5            required: ["app","uid","timestamp"],
6            properties: {
7                app: {
8                    bsonType: "string",
9                    description: "Application ID (string)"
10               },
11               uid: {
12                   bsonType: "string",
13                   description: "User (student) ID (string)"
14               },
15               context: {
16                   bsonType: "string",
17                   description: "Context (task) ID (string)"
18               },
19               sender: {
20                   bsonType: "string",
21                   description: "Who posted this message."
22               },
23               mess: {
24                   bsonType: "string",
25                   description: "Topic of Message"
26               },
27               timestamp: {
28                   bsonType: "date",
29                   description: "Timestamp"
30               },
31               data: {
32                   bsonType: "object",
33                   description: "Data about Activity Selection."
34               }
35           }
36       }
37   },
38   validationAction: "warn"
39 });
40 db.Activity.createIndex( { app:1, uid: 1, timestamp: -1});
```

PlayerStart.php Called when player logs in on a given day. As data returns information needed to restore gaming session (currently bank balance and list of trophies earned). Note that player details are updated by the EI process.

PlayerStop.php Called when player logs out. Currently not used. It is designed to help automatically shut down unneeded processed.

PlayerStats.php Called when current player competency estimates are required, e.g., when displaying player scores. It returns a list of statistics and their values in the data field; the exact statistics returned depend on the configuration of the EA process. This database collection is updated by the EA process after each game level is processed.

PlayerLevels.php Called when the game wants the next level. The message data should contain information about what topic the player is currently addressing and a list of played and unplayed levels, with the unplayed levels sorted so the next level according to protocol is first on the list. The complete list of levels should be returned so that if levels on the list have already been completed, a new level would be entered. Although the PHP script has been built, the AS process to feed it has not.

In addition to the primary files there are some auxiliary files that are available as well.

P4echo.php This script simply repeats back the message that was sent as a json object. Intended for testing.

Proc4.ini This contains configuration information used by the other processes, particularly, database credentials and a list of supported application names. A template file is provided in the `config` directory. It should be edited and moved to `/usr/local/share/Proc4` (or other designated configuration directory).

config.php This script is called by the others, it is mainly calls the `Proc4.ini` script, so this file can be modified if that script is in a non-standard location.

composer.json This file is generated by composer to get mongo to install.

The configuration file `Proc4.ini` contains a list of application IDs and the passwords for the databases. The file is shown in Listing 7. The applications are particularly important as they serve as a password for systems that use this facility. In particular, unless the `app` field of the POST request is one of the applications listed in the ini file, then the scripts will return an error. This should prevent random hacking, but more serious security might be needed if there is a more substantial risk.

Installation requires the following steps:

1. Edit the `Proc4.ini` file (in the `config` subdirectory of the `Proc4` package) and move it to the configuration directory, by default `/usr/local/share/Proc4`.

Listing 7: PHP initialization file Proc4.ini.

```
1  [apps]
2  test = "ecd://epls.coe.fsu.edu/P4test"
3  userControl = "ecd://epls.coe.fsu.edu/PhysicsPlayground/userControl"
4  linear = "ecd://epls.coe.fsu.edu/PhysicsPlayground/linear"
5  adaptive = "ecd://epls.coe.fsu.edu/PhysicsPlayground/adaptive"
6
7  [users]
8  EIP = "secret"
9  EAP = "secret"
10 ASP = "secret"
11 C4 = "secret"
```

2. Edit the `config.php` file if necessary and copy the php scripts to a directory exposed by the web server.

3. Install the mongo PHP drivers using PECL and composer. See the instructions at `https://docs.mongodb.com/ecosystem/drivers/php/`. (Note for RHEL. Because RHEL 7.5 is behind the curve on a large number of packages, the available drivers for RHEL have lower version numbers. You may need to remove the `composer.json` file. The dongle appears to work fine with version 1.1 of mongodb, which is what I get with RHEL 7.5.)

5 Pulling statements from the learning record store.

Learning Locker® stores events as `xAPI` (Betts & Smith, 2018) formatted JSON in a collection called `statements` in a database called `gameLRS` (or at least that is the setup for *Physics Playground*). All of the statements have a timestamp, so the extraction loop can get only new messages after the first extraction. The scripts `extractEvidence.sh` and `importEvidence.sh` facilitate the extraction from learning locker and the upload into the `EIRecords` database.

Between extraction and importation, the messages must be converted from xAPI (actually a wrapped xAPI format) to P4 format. This is done by the bash script `LLtoP4` (Listing 8). The translation is done in three steps. The first step, using the program `jq` (`https://stedolan.github.io/jq/`), extracts the fields relevant for the P4 messages from the unused information. Note that much of the useful information as defined in `Physics Playground` is in the extension for the object element of the statement. The second step uses standard GNU tool sed (Windows users, see `http://gnuwin32.sourceforge.net/packages/sed.htm`) to

Listing 8: LLtoP4 converter

```bash
#!/bin/bash
jq -f filter1.jq | sed -f filter2.sed | jq -f filter3.jq
```

shorten long URL-like guids to shorter keywords. The third step promotes some information (in particular, the `app` and `context` fields) which are in the extensions to the header.

In theory, simply looping the shell command, `extractEvidence.sh date | LLtoP4 | importEvidence.sh`, is all that is necessary. In practice, two additional steps are needed. First, it is necessary to extract the most recent timestamp from the downloaded file. The next extraction will be for all events after that timestamp. Second, it is often useful to filter the events before uploading them to the database.

Adding an extra filtering step to the extraction loop is a big time saver, as events which will not trigger any evidence rules can simply be discarded. The filter `jq -f coreEvents.jq` is used to delete events which will not trigger rules. The result was about a 500-fold reduction in the operational version of *Physics Playground*.

To allow for a graceful shutdown, the loop is given a name (the first argument to the shell script) and a file `/usr/local/share/Proc4/log/name.running` is created when the script starts. In every loop, that file is checked. If it no longer exists, the script exits. So the loop can be shut down by removing the file. (The script is typically started using `nohup LLtoP4Loop name date >../logs/name.log &` so that it runs as background process.) Listing 9 shows the listing.

6 Configuration

This manual assumes that a modify LAMP (Linux, Apache, Mongo, PHP) stack is configured on the target machine. (Sorry, Windows users, a fair amount of adaptation will be needed to run under Windows.) Refer to the help files for your Linux distribution and Mongo to accomplish this task. Also, for the Proc4 for R library, R (R Core Team, 2018) will need to be installed on the target system.

6.1 Configuration Files

The first step is to pick a configuration directory for Proc4. The current system assumes that the configuration directory is `/usr/local/share/Proc4/`, but this could be changed depending on local preferences. This directory will need to be created with root privileges, but can then be set as writable by a normal user account. Create two subdirectories `bin` and `logs` underneath the configuration directory. (The log directory could be a symlink to a directory on another partition if space on the root partition is at a premium). This directory

Listing 9: Download capture loop

```bash
#!/bin/bash
IP=127.0.0.1
name=$1
starttime=$2
echo "Learning Locker to P4 extraction loop, $1, starting: $2"
## Create a running file, when this file is deleted,
## the process will stop.
cd /usr/local/share/Proc4/bin
touch ../logs/$name.running

cache1=$(mktemp --tmpdir ${name}.XXXXXXXXX)
cache2=$(mktemp --tmpdir ${name}.XXXXXXXXX)

while [ -f ../logs/$name.running ]
do
    ssh $IP ./extractEvidence.sh $starttime >$cache1
    if [ $(tail -n +2 $cache1 | jq 'length') -gt 0 ]; then
        tail -n +2 <$cache1 | ./LLtoP4 | jq -f coreEvents.jq >$cache2
        ./importEvidence.sh <$cache2
        starttime=$(jq '[.[]|.timestamp."$date"]|max' $cache2)
    fi
    echo "Next extraction at $starttime"
    sleep 10s
done
```

will be called the `Proc4` directory in the sequel. Note that the location of the `Proc4` directory is hard coded into a number of files, and they will need to be manually edited if a different location is chosen.

Finally, there are a number of configuration files that are stored in the R package tarball. These can be accessed in one of two ways. First, install the R package, then use the command `library(help="Proc4")$path` to determine the install location for the R package. The `config` and `dongle` subdirectories are in that location. The alternative is to simply unpack the Proc4 tarball in some known location (equivalently, one could download from SVN in a given location). The subdirectories `inst/config` and `inst/dongle` contain the relevant files. These directories will be called the `config` and `dongle` directories respectively.

Copy the files `config/Proc4.js` and `config/Proc4.ini` to the `Proc4` directory. These files will need to be edited to reflect the local configuration. In particular, if database passwords are used, then they will need to be set in this file. Also, `Proc4.ini` has a list of valid applications. The EI, EA and AS processes will also store their initialization files in this directory.

6.2   Mongo Configuration

Using the mongo database, both security (user IDs and passwords) is optional. Running mongo without security turned on is probably okay as long as the installation is (a) behind a firewall, and (b) the firewall is configured to not allow connections on the mongo port except from localhost. However, other users may want to turn on security.

The recommended security setup is to create four users, "EIP", "EAP", "ASP", and "C4" for the four processes and to assign a password to each. The URI's of the database connections then need to be modified to include the username and passwords. Each process would have an `ini.R` file which contains its password which is stored in an appropriate configuration directory.

The files `Proc4/Proc4.ini` (PHP format) and `Proc4/Proc4.js` (javascript format) are used for saving the key usernames and passwords. Note that the mongo configuration files read the usernames and passwords from `Proc4/Proc4.js`, so this file needs to be configured before the following steps.

The file `setupDatabases.js` in the `config` directory creates databases for each of the processes and stores the appropriate login credentials. This is a javascript file designed to be run directly in mongo, i.e., `mongo setupDatabases.js`. Note that it must be run by a user which has the appropriate privileges to create databases and modify their security (a "root" user). This step is required if security is turned on in the database, and optional if it is turned off.

The file `config/setupProc4.js` sets up schemas and indexes for collections in the `Proc4` database which are used by the dongle process. Schemas are optional in mongo, but the indexes should speed up operations.

6.3   PHP Dongle Configuration

To create the dongle process, pick a directory under apache control (e.g., a subdirectory of `https_docs`) in which to install the Dongle. This will determine the URL base for the dongle scripts. Next, copy all of the PHP files and the file `composer.json` from the `dongle` directory to the web directory. If the `Proc4` directory is not at `/usr/local/share/Proc4`, then file `config.php` should be edited to reflect the proper path.

The file `Proc4/Proc4.ini` will need to be edited (a) to ensure the proper passwords are in place for the processes and (b) to list all of the legal applications in the `app` section. Note that the scripts use the `app` field to verify that the requester is actually associated with the project.

Ensure that the mongodb extensions for PHP have been installed (`https://docs.mongodb .com/ecosystem/drivers/php/`). Note that the last step is to run `composer` in the URL base directory for the dongle. (The supplied `composer.json` file was generated using Ubuntu 18.04. Under RHEL 7.5, an earlier version of the mongodb extension is needed. To install under RHEL, delete `composer.json` and install using `composer require mongodb/mongodb`.

The file `P4echo.php` can be used for testing the configuration. Simply point the browser at the file, and it will give you a form for sending a test message, which it will echo back. The other scripts work in a similar way, issuing a GET request (i.e., pointing a browser at the page) will return a form that can be used to POST a test message and return the JSON message. This may not be particularly useful until the databases have been populated though.

6.4   Event Loop Configuration

The remaining files (i.e., everything but the `php` files) in the `dongle` directory are for the event loop. These should be copied to `Proc4/bin`. Many of the shell scripts assume the location of the filter files (with `.jq` and `.sed` extensions) in the local directory, so the `LLtoP4Loop` command needs to be edited to run in this directory.

In the *Physics Playground* implementation learning locker and `Proc4` were run on different servers. To implement this, the IP address in `LLtoP4loop` needs to be updated to the name or IP address of the learning locker server. The file `extractEvidence.sh` also needs to be copied to the learning locker server and put in the login directory. The script uses an ssh tunnel to do the extraction; so this connection (both firewalls and ssh keys) needs to be properly configured.

Finally, the file `Proc4/bin/coreEvents.jq` determines which events are imported into the EI process queue. This will need editing depending on the rule set used by the EI process.

7 Acknowledgements

References

Almond, R. G., Steinberg, L. S., & Mislevy, R. J. (2002). Enhancing the design and delivery of assessment systems: A four-process architecture. *Journal of Technology, Learning, and Assessment*, *1*, (online). Retrieved from `http://www.jtla.org/`

Bassett, L. (2015). *Introduction to JavaScript object notation: A to-the-point guide to JSON.* O'Reilly Media, Inc.

Betts, B., & Smith, R. (2018). The leraning technology manager's guid to xAPI (2nd ed.) [Computer software manual]. Retrieved from `https://www.ht2labs.com/resources/the-learning-technology-managers-guide-to-the-xapi/#gf_26`

The mongodb 4.0 manual (4.0 ed.) [Computer software manual]. (2018). Retrieved from `https://docs.mongodb.com/manual/` (Retrieved 2018-09-03.)

R Core Team. (2018). R: A language and environment for statistical computing [Computer software manual]. Vienna, Austria. Retrieved from `https://www.R-project.org/`